

# CS 245: Database System Principles

## Notes 5: Hashing and More


Hector Garcia-Molina

CS 245

Notes 5

1

### Hashing




CS 245

Notes 5

2

### Two alternatives




CS 245

Notes 5

3

### Two alternatives




CS 245

Notes 5

4

### Two alternatives



- Alt (2) for "secondary" search key

CS 245

Notes 5

5

### Example hash function

- Key = 'x<sub>1</sub> x<sub>2</sub> ... x<sub>n</sub>' *n* byte character string
- Have *b* buckets
- h: add x<sub>1</sub> + x<sub>2</sub> + ... + x<sub>n</sub>
  - compute sum modulo *b*

CS 245

Notes 5

6

- ☒ This may not be best function ...
- ☒ Read Knuth Vol. 3 if you really need to select a good function.

CS 245

Notes 5

7

- ☒ This may not be best function ...
- ☒ Read Knuth Vol. 3 if you really need to select a good function.

Good hash function: ☞ Expected number of keys/bucket is the same for all buckets

CS 245

Notes 5

8

### Within a bucket:


- Do we keep keys sorted?
- Yes, if CPU time critical & Inserts/Deletes not too frequent

CS 245

Notes 5

9

Next: example to illustrate inserts, overflows, deletes




CS 245

Notes 5

10

### EXAMPLE 2 records/bucket




CS 245

Notes 5

11

### EXAMPLE 2 records/bucket



CS 245

Notes 5

12

### EXAMPLE 2 records/bucket

INSERT:


$h(a) = 1$

$h(b) = 2$

$h(c) = 1$

$h(d) = 0$

$h(e) = 1$



CS 245

Notes 5


13

### EXAMPLE: deletion

Delete:

e

f



CS 245

Notes 5

14


### EXAMPLE: deletion

Delete:

e

f

c



CS 245

Notes 5

15


### EXAMPLE: deletion

Delete:

e

f

c



CS 245

Notes 5

16

### Rule of thumb:

- Try to keep space utilization between 50% and 80%

$$\text{Utilization} = \frac{\# \text{ keys used}}{\text{total } \# \text{ keys that fit}}$$

CS 245

Notes 5

17

### Rule of thumb:

- Try to keep space utilization between 50% and 80%

$$\text{Utilization} = \frac{\# \text{ keys used}}{\text{total } \# \text{ keys that fit}}$$

- If < 50%, wasting space
- If > 80%, overflows significant
  - ↳ depends on how good hash function is & on # keys/bucket

CS 245

Notes 5

18

### How do we cope with growth?

- Overflows and reorganizations
- Dynamic hashing

CS 245

Notes 5

19

### How do we cope with growth?

- Overflows and reorganizations
- Dynamic hashing
  - Extensible
  - Linear


CS 245

Notes 5

20

### Extensible hashing: two ideas

(a) Use  $i$  of  $b$  bits output by hash function




CS 245

Notes 5

21

(b) Use directory




CS 245

Notes 5

22

### Example: $h(k)$ is 4 bits; 2 keys/bucket




Insert 1010

CS 245

Notes 5

23

### Example: $h(k)$ is 4 bits; 2 keys/bucket




Insert 1010

CS 245

Notes 5

24

### Example: $h(k)$ is 4 bits; 2 keys/bucket




CS 245

Notes 5

25

### Example continued

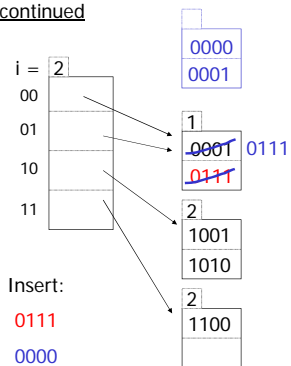


CS 245

Notes 5

26

### Example continued




CS 245

Notes 5

27

### Example continued




CS 245

Notes 5

28

### Example continued




CS 245

Notes 5

29

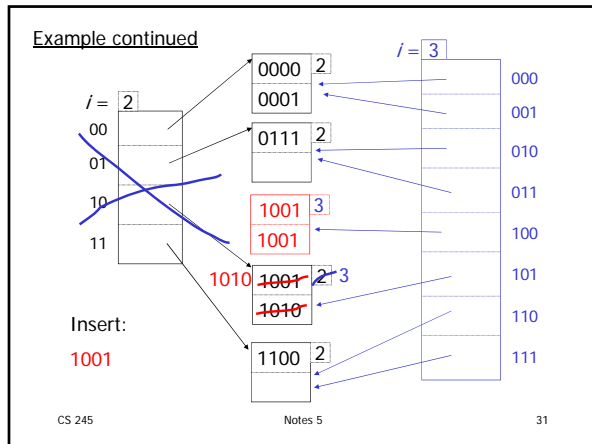
### Example continued



CS 245

Notes 5

30



### Extensible hashing: deletion

- No merging of blocks
- Merge blocks and cut directory if possible (Reverse insert procedure)

CS 245 Notes 5 32

### Deletion example:

- Run thru insert example in reverse!

CS 245 Notes 5 33

### Note: Still need overflow chains

- Example: many records with duplicate keys

insert 1100 if we split:

CS 245 Notes 5 34

### Solution: overflow chains

insert 1100 add overflow block:

CS 245 Notes 5 35

### Summary Extensible hashing

- ⊕ Can handle growing files
  - with less wasted space
  - with no full reorganizations

CS 245 Notes 5 36

## Summary Extensible hashing

- ⊕ Can handle growing files
  - with less wasted space
  - with no full reorganizations
- ⊖ Indirection  
(Not bad if directory in memory)
- ⊖ Directory doubles in size  
(Now it fits, now it does not)

CS 245

Notes 5


37

## Linear hashing

- Another dynamic hashing scheme

### Two ideas:

- (a) Use  $i$  low order bits of hash



CS 245

Notes 5


38

## Linear hashing

- Another dynamic hashing scheme

### Two ideas:

- (a) Use  $i$  low order bits of hash



- (b) File grows linearly




CS 245

Notes 5

39

## Example $b=4$ bits, $i=2$ , 2 keys/bucket




CS 245

Notes 5

40

## Example $b=4$ bits, $i=2$ , 2 keys/bucket



**Rule** If  $h(k)[i] \leq m$ , then  
 look at bucket  $h(k)[i]$   
 else, look at bucket  $h(k)[i] - 2^{i-1}$


CS 245

Notes 5

41

## Example $b=4$ bits, $i=2$ , 2 keys/bucket

- insert 0101



**Rule** If  $h(k)[i] \leq m$ , then  
 look at bucket  $h(k)[i]$   
 else, look at bucket  $h(k)[i] - 2^{i-1}$

CS 245

Notes 5

42

• insert 0101


• can have overflow chains!

Future growth buckets

$m = 01$  (max used block)



Example  $b=4$  bits,  $i=2$ , 2 keys/bucket



CS 245

Notes 5

49

Example Continued: How to grow beyond this?

$i = 2$

0000	0101	1010	1111
	0101		

$m = 11$  (max used block)

CS 245

Notes 5

50

Example Continued: How to grow beyond this?

$i = 2 \rightarrow 3$

0000	0101	1010	1111		
	0101				

000 001 010 011 ...  
100 101 110 111

$m = 11$  (max used block)

CS 245

Notes 5

51

Example Continued: How to grow beyond this?

$i = 2 \rightarrow 3$

0000	0101	1010	1111		
	0101				

000 001 010 011 100 ...  
100 101 110 111

$m = 11$  (max used block)

100

CS 245

Notes 5

52

Example Continued: How to grow beyond this?

$i = 2 \rightarrow 3$

0000	0101	1010	1111		0101
	0101				0101

000 001 010 011 100 101 ...  
100 101 110 111

$m = 11$  (max used block)

100

101

CS 245

Notes 5

53

☒ When do we expand file?

- Keep track of:  $\frac{\text{\# used slots}}{\text{total \# of slots}} = U$

CS 245

Notes 5

54

### ☒ When do we expand file?

- Keep track of:  $\frac{\text{\# used slots}}{\text{total \# of slots}} = U$
- If  $U > \text{threshold}$  then increase  $m$   
(and maybe  $i$ )

CS 245

Notes 5

55

### Summary Linear Hashing


- ⊕ Can handle growing files
  - with less wasted space
  - with no full reorganizations
- ⊕ No indirection like extensible hashing
- ⊖ Can still have overflow chains

CS 245

Notes 5

56

### Example: BAD CASE



CS 245

Notes 5

57

### Summary

#### Hashing

- How it works
- Dynamic hashing
  - Extensible
  - Linear

CS 245

Notes 5

58

### Next:

- Indexing vs Hashing
- Index definition in SQL
- Multiple key access

CS 245

Notes 5

59

### Indexing vs Hashing

- Hashing good for probes given key  
e.g.,  
SELECT ...  
FROM R  
WHERE R.A = 5

CS 245

Notes 5

60

## Indexing vs Hashing

- INDEXING (Including B Trees) good for Range Searches:  
e.g.,  
SELECT  
FROM R  
WHERE R.A > 5

CS 245

Notes 5

61

## Index definition in SQL

- Create index name on rel (attr)
- Create unique index name on rel (attr)  
└─→ defines candidate key
- Drop INDEX name

CS 245

Notes 5

62

**Note** CANNOT SPECIFY TYPE OF INDEX  
(e.g. B-tree, Hashing, ...)  
OR PARAMETERS  
(e.g. Load Factor, Size of Hash,...)

... at least in SQL...

CS 245

Notes 5

63

**Note** ATTRIBUTE LIST ⇒ MULTIKEY INDEX  
(next)  
e.g., CREATE INDEX foo ON R(A,B,C)

CS 245

Notes 5

64

## Multi-key Index

Motivation: Find records where  
DEPT = "Toy" AND SAL > 50k


CS 245

Notes 5

65

## Strategy I:

- Use one index, say Dept.
- Get all Dept = "Toy" records  
and check their salary



CS 245

Notes 5

66

### Strategy II:

- Use 2 Indexes; Manipulate Pointers

Toy → 

--	--	--	--	--	--	--	--	--	--

--	--	--	--	--	--	--	--	--	--

 ← Sal > 50k

CS 245


Notes 5

67

### Strategy III:

- Multiple Key Index

One idea:




CS 245

Notes 5

68

### Example



CS 245

Notes 5

69

For which queries is this index good?

- ☐ Find RECs Dept = "Sales"  $\wedge$  SAL=20k
- ☐ Find RECs Dept = "Sales"  $\wedge$  SAL  $\geq$  20k
- ☐ Find RECs Dept = "Sales"
- ☐ Find RECs SAL = 20k


CS 245

Notes 5

70

### Interesting application:

- Geographic Data



DATA:  
<X1,Y1, Attributes>  
<X2,Y2, Attributes>  
⋮

CS 245

Notes 5

71

### Queries:


- What city is at <Xi,Yi>?
- What is within 5 miles from <Xi,Yi>?
- Which is closest point to <Xi,Yi>?

CS 245

Notes 5

72

### Example





CS 245

Notes 5

73

### Example




CS 245

Notes 5

74

### Example

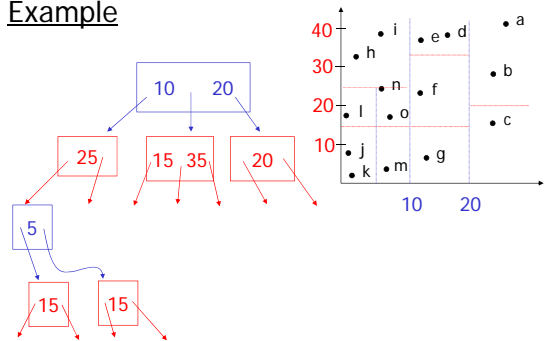


CS 245

Notes 5

75

### Example

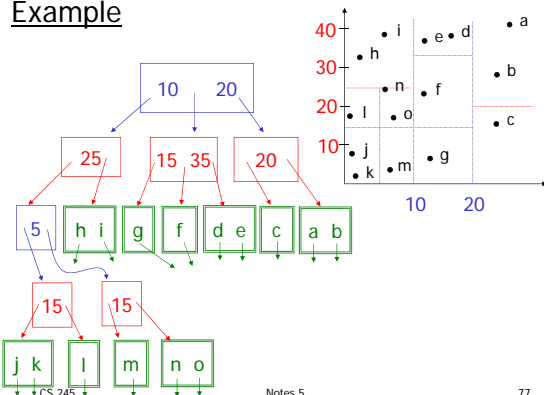


CS 245

Notes 5

76

### Example

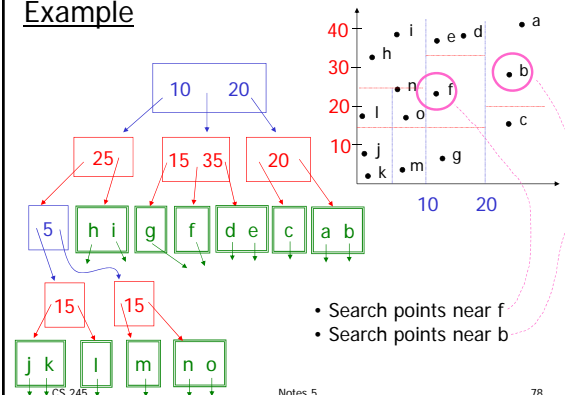


CS 245

Notes 5

77

### Example



Notes 5

78

### Queries

- Find points with  $Y_i > 20$
- Find points with  $X_i < 5$
- Find points "close" to  $i = \langle 12, 38 \rangle$
- Find points "close" to  $b = \langle 7, 24 \rangle$

CS 245

Notes 5

79

- Many types of geographic index structures have been suggested
  - kd-Trees (very similar to what we described here)
  - Quad Trees
  - R Trees
  - ...

CS 245

Notes 5

80

### Two more types of multi key indexes


- Grid
- Partitioned hash

CS 245

Notes 5

81

### Grid Index



CS 245

Notes 5

82

### CLAIM

- Can quickly find records with
  - key 1 =  $V_i \wedge$  Key 2 =  $X_j$
  - key 1 =  $V_i$
  - key 2 =  $X_j$

CS 245

Notes 5

83

### CLAIM


- Can quickly find records with
  - key 1 =  $V_i \wedge$  Key 2 =  $X_j$
  - key 1 =  $V_i$
  - key 2 =  $X_j$
- And also ranges....
  - E.g., key 1  $\geq V_i \wedge$  key 2  $< X_j$

CS 245

Notes 5

84

- How do we find entry  $i, j$  in linear structure?




CS 245

Notes 5

85

- How do we find entry  $i, j$  in linear structure?




CS 245

Notes 5

86

### Solution: Use Indirection



CS 245

Notes 5

87

### With indirection:


- Grid can be regular without wasting space
- We do have price of indirection

CS 245

Notes 5

88

### Can also index grid on value ranges



CS 245

Notes 5

89

### Grid files

- ⊕ Good for multiple-key search
- ⊖ Space, management overhead (nothing is free)
- ⊖ Need partitioning ranges that evenly split keys


CS 245

Notes 5

90

## Partitioned hash function

Idea:



CS 245

Notes 5

91

EX:

h1(toy)	=0	000	
h1(sales)	=1	001	
h1(art)	=1	010	
.		011	
h2(10k)	=01	100	
h2(20k)	=11	101	
h2(30k)	=01	110	
h2(40k)	=00	111	
.			



<Fred,toy,10k>, <Joe,sales,10k>  
<Sally,art,30k>

CS 245

Notes 5

92

EX:

h1(toy)	=0	000	
h1(sales)	=1	001	<Fred>
h1(art)	=1	010	
.		011	
h2(10k)	=01	100	
h2(20k)	=11	101	<Joe><Sally>
h2(30k)	=01	110	
h2(40k)	=00	111	
.			



<Fred,toy,10k>, <Joe,sales,10k>  
<Sally,art,30k>

CS 245

Notes 5

93

h1(toy)	=0	000	<Fred>
h1(sales)	=1	001	<Joe><Jan>
h1(art)	=1	010	<Mary>
.		011	
h2(10k)	=01	100	<Sally>
h2(20k)	=11	101	
h2(30k)	=01	110	<Tom><Bill>
h2(40k)	=00	111	<Andy>
.			

- Find Emp. with Dept. = Sales  $\wedge$  Sal=40k

CS 245

Notes 5

94

h1(toy)	=0	000	<Fred>
h1(sales)	=1	001	<Joe><Jan>
h1(art)	=1	010	<Mary>
.		011	
h2(10k)	=01	100	<Sally>
h2(20k)	=11	101	
h2(30k)	=01	110	<Tom><Bill>
h2(40k)	=00	111	<Andy>
.			

- Find Emp. with Dept. = Sales  $\wedge$  Sal=40k

CS 245

Notes 5

95

h1(toy)	=0	000	<Fred>
h1(sales)	=1	001	<Joe><Jan>
h1(art)	=1	010	<Mary>
.		011	
h2(10k)	=01	100	<Sally>
h2(20k)	=11	101	
h2(30k)	=01	110	<Tom><Bill>
h2(40k)	=00	111	<Andy>
.			

- Find Emp. with Sal=30k

CS 245

Notes 5

96



h1(toy)	=0	000	<Fred>
h1(sales)	=1	001	<Joe><Jan>
h1(art)	=1	010	<Mary>
.	.	011	
h2(10k)	=01	100	<Sally>
h2(20k)	=11	101	
h2(30k)	=01	110	<Tom><Bill>
h2(40k)	=00	111	<Andy>
.	.		

• Find Emp. with Sal=30k

look here

CS 245

Notes 5

97

h1(toy)	=0	000	<Fred>
h1(sales)	=1	001	<Joe><Jan>
h1(art)	=1	010	<Mary>
.	.	011	
h2(10k)	=01	100	<Sally>
h2(20k)	=11	101	
h2(30k)	=01	110	<Tom><Bill>
h2(40k)	=00	111	<Andy>
.	.		

• Find Emp. with Dept. = Sales

CS 245

Notes 5

98

h1(toy)	=0	000	<Fred>
h1(sales)	=1	001	<Joe><Jan>
h1(art)	=1	010	<Mary>
.	.	011	
h2(10k)	=01	100	<Sally>
h2(20k)	=11	101	
h2(30k)	=01	110	<Tom><Bill>
h2(40k)	=00	111	<Andy>
.	.		

• Find Emp. with Dept. = Sales

look here

CS 245

Notes 5

99

## Summary

### Post hashing discussion:

- Indexing vs. Hashing
- SQL Index Definition
- Multiple Key Access
  - Multi Key Index
    - Variations: Grid, Geo Data
- Partitioned Hash

CS 245

Notes 5

100

## Reading Chapter 5

- Skim the following sections:
  - Sections 14.3.6, 14.3.7, 14.3.8  
[Second Ed: 14.6.6, 14.6.7, 14.6.8]
  - Sections 14.4.2, 14.4.3, 14.4.4  
[Second Ed: 14.7.2, 14.7.3, 14.7.4]
- Read the rest

CS 245

Notes 5

101

## The BIG picture....

- Chapters 11 & 12 [13]: Storage, records, blocks...
- Chapters 13 & 14 [14]: Access Mechanisms
  - Indexes
  - B trees
  - Hashing
  - Multi key
- Chapters 15 & 16 [15, 16]: Query Processing

NEXT →

CS 245

Notes 5

102